

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
In [2]: import warnings

warnings.filterwarnings('ignore')
```

```
In [3]: data = 'C:/Users/surya/Downloads/sistem-cerdas/adult.csv'

df = pd.read_csv(data, header=None, sep=',\s')
```

```
In [4]: df.shape
```

```
Out[4]: (32561, 15)
```

```
In [5]: df.head()
```

```
Out[5]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40

```
In [6]: col_names = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status',
                    'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country']

df.columns = col_names

df.columns
```

```
Out[6]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
               'marital_status', 'occupation', 'relationship', 'race', 'sex',
```

```
'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
'income'],
dtype='object')
```

In [7]: `df.head()`

Out[7]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education_num          32561 non-null  int64
5   marital_status         32561 non-null  object
6   occupation             32561 non-null  object
7   relationship           32561 non-null  object
8   race                   32561 non-null  object
9   sex                    32561 non-null  object
10  capital_gain           32561 non-null  int64
11  capital_loss           32561 non-null  int64
12  hours_per_week         32561 non-null  int64
13  native_country         32561 non-null  object
14  income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

In [9]:

```
categorical = [var for var in df.columns if df[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are :\n\n', categorical)
```

There are 9 categorical variables

The categorical variables are :

```
['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race',
'sex', 'native_country', 'income']
```

In [10]: `df[categorical].head()`

Out[10]:

	workclass	education	marital_status	occupation	relationship	race	sex	native_country	inc
0	State-gov	Bachelors	Never-married	Adm-clerical	Not-in-family	White	Male	United-States	<
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States	<
2	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States	<
3	Private	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	United-States	<
4	Private	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	Cuba	<



In [11]:

```
df[categorical].isnull().sum()
```

Out[11]:

```
workclass      0
education      0
marital_status 0
occupation     0
relationship   0
race           0
sex            0
native_country 0
income         0
dtype: int64
```

In [12]:

```
for var in categorical:
    print(df[var].value_counts())
```

```
Private      22696
Self-emp-not-inc  2541
Local-gov    2093
?            1836
State-gov    1298
Self-emp-inc  1116
Federal-gov   960
Without-pay   14
Never-worked   7
Name: workclass, dtype: int64
HS-grad      10501
Some-college  7291
Bachelors    5355
Masters       1723
Assoc-voc    1382
11th         1175
Assoc-acdm   1067
10th         933
7th-8th      646
Prof-school   576
9th          514
12th         433
Doctorate     413
5th-6th       333
1st-4th       168
Preschool     51
Name: education, dtype: int64
Married-civ-spouse  14976
Never-married      10683
Divorced            4443
```

```

Separated                1025
Widowed                  993
Married-spouse-absent    418
Married-AF-spouse        23
Name: marital_status, dtype: int64
Prof-specialty           4140
Craft-repair             4099
Exec-managerial          4066
Adm-clerical             3770
Sales                   3650
Other-service            3295
Machine-op-inspct        2002
?                        1843
Transport-moving         1597
Handlers-cleaners        1370
Farming-fishing          994
Tech-support             928
Protective-serv          649
Priv-house-serv          149
Armed-Forces             9
Name: occupation, dtype: int64
Husband                 13193
Not-in-family           8305
Own-child               5068
Unmarried               3446
Wife                   1568
Other-relative          981
Name: relationship, dtype: int64
White                   27816
Black                   3124
Asian-Pac-Islander      1039
Amer-Indian-Eskimo      311
Other                   271
Name: race, dtype: int64
Male                   21790
Female                 10771
Name: sex, dtype: int64
United-States           29170
Mexico                  643
?                       583
Philippines             198
Germany                 137
Canada                 121
Puerto-Rico            114
El-Salvador            106
India                  100
Cuba                   95
England                90
Jamaica                81
South                  80
China                  75
Italy                  73
Dominican-Republic     70
Vietnam                67
Guatemala              64
Japan                  62
Poland                 60
Columbia               59
Taiwan                 51
Haiti                  44
Iran                   43
Portugal               37
Nicaragua              34
Peru                   31
Greece                 29
France                 29
Ecuador                28
Ireland                24
Hong                   20

```

```

Cambodia                19
Trinidad&Tobago         19
Laos                    18
Thailand                 18
Yugoslavia              16
Outlying-US(Guam-USVI-etc) 14
Hungary                 13
Honduras                13
Scotland                12
Holand-Netherlands      1
Name: native_country, dtype: int64
<=50K      24720
>50K        7841
Name: income, dtype: int64

```

In [13]:

```

for var in categorical:

    print(df[var].value_counts()/np.float(len(df)))

```

```

Private                0.697030
Self-emp-not-inc      0.078038
Local-gov             0.064279
?                     0.056386
State-gov             0.039864
Self-emp-inc          0.034274
Federal-gov           0.029483
Without-pay           0.000430
Never-worked          0.000215
Name: workclass, dtype: float64
HS-grad              0.322502
Some-college         0.223918
Bachelors            0.164461
Masters              0.052916
Assoc-voc            0.042443
11th                 0.036086
Assoc-acdm           0.032769
10th                 0.028654
7th-8th              0.019840
Prof-school          0.017690
9th                  0.015786
12th                 0.013298
Doctorate            0.012684
5th-6th              0.010227
1st-4th              0.005160
Preschool            0.001566
Name: education, dtype: float64
Married-civ-spouse    0.459937
Never-married         0.328092
Divorced              0.136452
Separated             0.031479
Widowed              0.030497
Married-spouse-absent 0.012837
Married-AF-spouse     0.000706
Name: marital_status, dtype: float64
Prof-specialty        0.127146
Craft-repair          0.125887
Exec-managerial       0.124873
Adm-clerical          0.115783
Sales                 0.112097
Other-service         0.101195
Machine-op-inspct     0.061485
?                     0.056601
Transport-moving      0.049046
Handlers-cleaners     0.042075
Farming-fishing       0.030527
Tech-support          0.028500
Protective-serv       0.019932
Priv-house-serv       0.004576

```

```

Armed-Forces      0.000276
Name: occupation, dtype: float64
Husband          0.405178
Not-in-family    0.255060
Own-child        0.155646
Unmarried        0.105832
Wife             0.048156
Other-relative   0.030128
Name: relationship, dtype: float64
White            0.854274
Black           0.095943
Asian-Pac-Islander 0.031909
Amer-Indian-Eskimo 0.009551
Other            0.008323
Name: race, dtype: float64
Male            0.669205
Female          0.330795
Name: sex, dtype: float64
United-States    0.895857
Mexico          0.019748
?               0.017905
Philippines     0.006081
Germany         0.004207
Canada          0.003716
Puerto-Rico     0.003501
El-Salvador     0.003255
India           0.003071
Cuba            0.002918
England         0.002764
Jamaica         0.002488
South           0.002457
China           0.002303
Italy           0.002242
Dominican-Republic 0.002150
Vietnam         0.002058
Guatemala       0.001966
Japan           0.001904
Poland          0.001843
Columbia        0.001812
Taiwan          0.001566
Haiti           0.001351
Iran            0.001321
Portugal        0.001136
Nicaragua       0.001044
Peru            0.000952
Greece          0.000891
France          0.000891
Ecuador         0.000860
Ireland         0.000737
Hong            0.000614
Cambodia        0.000584
Trinidad&Tobago 0.000584
Laos            0.000553
Thailand         0.000553
Yugoslavia      0.000491
Outlying-US(Guam-USVI-etc) 0.000430
Hungary         0.000399
Honduras        0.000399
Scotland        0.000369
Holand-Netherlands 0.000031
Name: native_country, dtype: float64
<=50K          0.75919
>50K           0.24081
Name: income, dtype: float64

```

```
In [14]: df.workclass.unique()
```

```
Out[14]: array(['State-gov', 'Self-emp-not-inc', 'Private', 'Federal-gov',
```

```
'Local-gov', '?', 'Self-emp-inc', 'Without-pay', 'Never-worked'],
dtype=object)
```

```
In [15]: df.workclass.value_counts()
```

```
Out[15]: Private          22696
Self-emp-not-inc      2541
Local-gov             2093
?                     1836
State-gov             1298
Self-emp-inc          1116
Federal-gov           960
Without-pay           14
Never-worked           7
Name: workclass, dtype: int64
```

```
In [16]: df['workclass'].replace('?', np.NaN, inplace=True)
```

```
In [17]: df.workclass.value_counts()
```

```
Out[17]: Private          22696
Self-emp-not-inc      2541
Local-gov             2093
State-gov             1298
Self-emp-inc          1116
Federal-gov           960
Without-pay           14
Never-worked           7
Name: workclass, dtype: int64
```

```
In [18]: df.occupation.unique()
```

```
Out[18]: array(['Adm-clerical', 'Exec-managerial', 'Handlers-cleaners',
'Prof-specialty', 'Other-service', 'Sales', 'Craft-repair',
'Transport-moving', 'Farming-fishing', 'Machine-op-inspct',
'Tech-support', '?', 'Protective-serv', 'Armed-Forces',
'Priv-house-serv'], dtype=object)
```

```
In [19]: df.occupation.value_counts()
```

```
Out[19]: Prof-specialty      4140
Craft-repair      4099
Exec-managerial    4066
Adm-clerical      3770
Sales             3650
Other-service     3295
Machine-op-inspct 2002
?                1843
Transport-moving  1597
Handlers-cleaners 1370
Farming-fishing   994
Tech-support      928
Protective-serv   649
Priv-house-serv   149
Armed-Forces       9
Name: occupation, dtype: int64
```

```
In [20]: df['occupation'].replace('?', np.NaN, inplace=True)
```

```
In [21]: df.occupation.value_counts()
```

```
Out[21]: Prof-specialty      4140
         Craft-repair       4099
         Exec-managerial    4066
         Adm-clerical       3770
         Sales              3650
         Other-service      3295
         Machine-op-inspct  2002
         Transport-moving   1597
         Handlers-cleaners  1370
         Farming-fishing    994
         Tech-support       928
         Protective-serv    649
         Priv-house-serv    149
         Armed-Forces       9
         Name: occupation, dtype: int64
```

```
In [22]: df.native_country.unique()
```

```
Out[22]: array(['United-States', 'Cuba', 'Jamaica', 'India', '?', 'Mexico',
                'South', 'Puerto-Rico', 'Honduras', 'England', 'Canada', 'Germany',
                'Iran', 'Philippines', 'Italy', 'Poland', 'Columbia', 'Cambodia',
                'Thailand', 'Ecuador', 'Laos', 'Taiwan', 'Haiti', 'Portugal',
                'Dominican-Republic', 'El-Salvador', 'France', 'Guatemala',
                'China', 'Japan', 'Yugoslavia', 'Peru',
                'Outlying-US(Guam-USVI-etc)', 'Scotland', 'Trinidad&Tobago',
                'Greece', 'Nicaragua', 'Vietnam', 'Hong', 'Ireland', 'Hungary',
                'Holand-Netherlands'], dtype=object)
```

```
In [23]: df.native_country.value_counts()
```

```
Out[23]: United-States      29170
         Mexico             643
         ?                  583
         Philippines        198
         Germany            137
         Canada             121
         Puerto-Rico        114
         El-Salvador        106
         India              100
         Cuba               95
         England            90
         Jamaica            81
         South              80
         China              75
         Italy              73
         Dominican-Republic 70
         Vietnam            67
         Guatemala         64
         Japan              62
         Poland             60
         Columbia          59
         Taiwan             51
         Haiti              44
         Iran               43
         Portugal           37
         Nicaragua          34
         Peru               31
         Greece             29
         France             29
         Ecuador            28
         Ireland            24
         Hong               20
         Cambodia           19
         Trinidad&Tobago    19
         Laos               18
         Thailand           18
```



```
Yugoslavia          16
Outlying-US(Guam-USVI-etc) 14
Hungary             13
Honduras            13
Scotland            12
Holand-Netherlands  1
Name: native_country, dtype: int64
```

```
In [24]: df['native_country'].replace('?', np.NaN, inplace=True)
```

```
In [25]: df.native_country.value_counts()
```

```
Out[25]: United-States      29170
Mexico                    643
Philippines              198
Germany                  137
Canada                   121
Puerto-Rico             114
El-Salvador              106
India                    100
Cuba                      95
England                  90
Jamaica                   81
South                     80
China                     75
Italy                     73
Dominican-Republic       70
Vietnam                   67
Guatemala                 64
Japan                     62
Poland                    60
Columbia                  59
Taiwan                    51
Haiti                     44
Iran                      43
Portugal                  37
Nicaragua                 34
Peru                      31
Greece                    29
France                    29
Ecuador                   28
Ireland                   24
Hong                      20
Trinidad&Tobago           19
Cambodia                  19
Thailand                  18
Laos                      18
Yugoslavia                16
Outlying-US(Guam-USVI-etc) 14
Hungary                   13
Honduras                  13
Scotland                  12
Holand-Netherlands        1
Name: native_country, dtype: int64
```

```
In [26]: df[categorical].isnull().sum()
```

```
Out[26]: workclass      1836
education      0
marital_status  0
occupation     1843
relationship   0
race           0
sex            0
native_country  583
```

```
income          0
dtype: int64
```

```
In [27]: for var in categorical:

        print(var, ' contains ', len(df[var].unique()), ' labels')
```

```
workclass contains 9 labels
education contains 16 labels
marital_status contains 7 labels
occupation contains 15 labels
relationship contains 6 labels
race contains 5 labels
sex contains 2 labels
native_country contains 42 labels
income contains 2 labels
```

```
In [28]: numerical = [var for var in df.columns if df[var].dtype!='O']

        print('There are {} numerical variables\n'.format(len(numerical)))

        print('The numerical variables are :', numerical)
```

There are 6 numerical variables

The numerical variables are : ['age', 'fnlwgt', 'education_num', 'capital_gain', 'capital_loss', 'hours_per_week']

```
In [29]: df[numerical].head()
```

```
Out[29]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
0	39	77516	13	2174	0	40
1	50	83311	13	0	0	13
2	38	215646	9	0	0	40
3	53	234721	7	0	0	40
4	28	338409	13	0	0	40

```
In [30]: df[numerical].isnull().sum()
```

```
Out[30]: age          0
fnlwgt        0
education_num  0
capital_gain   0
capital_loss   0
hours_per_week 0
dtype: int64
```

```
In [31]: X = df.drop(['income'], axis=1)

        y = df['income']
```

```
In [32]: from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_st
```

```
In [33]: X_train.shape, X_test.shape
```

```
Out[33]: ((22792, 14), (9769, 14))
```

```
In [34]: X_train.dtypes
```

```
Out[34]: age                int64
workclass                 object
fnlwgt                   int64
education                 object
education_num            int64
marital_status           object
occupation                object
relationship              object
race                     object
sex                      object
capital_gain             int64
capital_loss             int64
hours_per_week           int64
native_country           object
dtype: object
```

```
In [35]: categorical = [col for col in X_train.columns if X_train[col].dtypes == 'O']

categorical
```

```
Out[35]: ['workclass',
'education',
'marital_status',
'occupation',
'relationship',
'race',
'sex',
'native_country']
```

```
In [36]: numerical = [col for col in X_train.columns if X_train[col].dtypes != 'O']

numerical
```

```
Out[36]: ['age',
'fnlwgt',
'education_num',
'capital_gain',
'capital_loss',
'hours_per_week']
```

```
In [37]: X_train[categorical].isnull().mean()
```

```
Out[37]: workclass        0.055985
education      0.000000
marital_status 0.000000
occupation     0.056072
relationship   0.000000
race           0.000000
sex            0.000000
native_country 0.018164
dtype: float64
```

```
In [38]: for col in categorical:
          if X_train[col].isnull().mean()>0:
              print(col, (X_train[col].isnull().mean()))
```

```
workclass 0.055984555984555984
occupation 0.05607230607230607
native_country 0.018164268164268166
```

```
In [39]: for df2 in [X_train, X_test]:
          df2['workclass'].fillna(X_train['workclass'].mode()[0], inplace=True)
          df2['occupation'].fillna(X_train['occupation'].mode()[0], inplace=True)
          df2['native_country'].fillna(X_train['native_country'].mode()[0], inplace=True)
```

```
In [40]: X_train[categorical].isnull().sum()
```

```
Out[40]: workclass      0
          education     0
          marital_status 0
          occupation     0
          relationship    0
          race           0
          sex            0
          native_country  0
          dtype: int64
```

```
In [41]: X_test[categorical].isnull().sum()
```

```
Out[41]: workclass      0
          education     0
          marital_status 0
          occupation     0
          relationship    0
          race           0
          sex            0
          native_country  0
          dtype: int64
```

```
In [42]: X_train.isnull().sum()
```

```
Out[42]: age           0
          workclass     0
          fnlwgt        0
          education     0
          education_num  0
          marital_status 0
          occupation     0
          relationship    0
          race           0
          sex            0
          capital_gain    0
          capital_loss    0
          hours_per_week  0
          native_country  0
          dtype: int64
```

```
In [43]: X_test.isnull().sum()
```

```
Out[43]: age           0
          workclass     0
          fnlwgt        0
          education     0
          education_num  0
          marital_status 0
          occupation     0
          relationship    0
          race           0
```

```
sex          0
capital_gain 0
capital_loss 0
hours_per_week 0
native_country 0
dtype: int64
```

In [44]: categorical

Out[44]: ['workclass',
'education',
'marital_status',
'occupation',
'relationship',
'race',
'sex',
'native_country']

In [45]: X_train[categorical].head()

Out[45]:

	workclass	education	marital_status	occupation	relationship	race	sex	native_country
32098	Private	HS-grad	Married-civ-spouse	Craft-repair	Husband	White	Male	United-States
25206	State-gov	HS-grad	Divorced	Adm-clerical	Unmarried	White	Female	United-States
23491	Private	Some-college	Married-civ-spouse	Sales	Husband	White	Male	United-States
12367	Private	HS-grad	Never-married	Craft-repair	Not-in-family	White	Male	Guatemala
7054	Private	7th-8th	Never-married	Craft-repair	Not-in-family	White	Male	Germany

In [48]: import category_encoders as ce

In [49]: encoder = ce.OneHotEncoder(cols=['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country'])

X_train = encoder.fit_transform(X_train)


X_test = encoder.transform(X_test)

In [50]: X_train.head()

Out[50]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7
32098	45	1	0	0	0	0	0	0
25206	47	0	1	0	0	0	0	0
23491	48	1	0	0	0	0	0	0
12367	29	1	0	0	0	0	0	0
7054	23	1	0	0	0	0	0	0

5 rows × 105 columns



```
In [51]: X_train.shape
```

Out[51]: (22792, 105)

```
In [52]: X_test.head()
```

Out[52]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7
22278	27	1	0	0	0	0	0	0
8950	27	1	0	0	0	0	0	0
7838	25	1	0	0	0	0	0	0
16505	46	1	0	0	0	0	0	0
19140	45	1	0	0	0	0	0	0

5 rows × 105 columns



```
In [53]: X_test.shape
```

Out[53]: (9769, 105)

```
In [54]: cols = X_train.columns
```

```
In [55]: from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)
```

```
In [56]: X_train = pd.DataFrame(X_train, columns=cols)
```

```
In [57]: X_test = pd.DataFrame(X_test, columns=cols)
```


```
In [58]: X_train.head()
```

Out[58]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7
0	0.40	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.50	-1.0	1.0	0.0	0.0	0.0	0.0	0.0
2	0.55	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	-0.40	0.0	0.0	0.0	0.0	0.0	0.0	0.0

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7
4	-0.70	0.0	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 105 columns



```
In [59]: from sklearn.naive_bayes import GaussianNB

gnb = GaussianNB()

gnb.fit(X_train, y_train)
```

Out[59]: GaussianNB()

```
In [60]: y_pred = gnb.predict(X_test)

y_pred
```

Out[60]: array(['<=50K', '<=50K', '>50K', ..., '>50K', '<=50K', '<=50K'],
dtype='<U5')

```
In [61]: from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy score: 0.8083

```
In [62]: y_pred_train = gnb.predict(X_train)

y_pred_train
```

Out[62]: array(['>50K', '<=50K', '>50K', ..., '<=50K', '>50K', '<=50K'],
dtype='<U5')

```
In [63]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

Training-set accuracy score: 0.8067

```
In [64]: print('Training set score: {:.4f}'.format(gnb.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(gnb.score(X_test, y_test)))
```

Training set score: 0.8067
Test set score: 0.8083

```
In [65]: y_test.value_counts()
```

Out[65]: <=50K 7407
>50K 2362
Name: income, dtype: int64

```
In [66]: null_accuracy = (7407/(7407+2362))

print('Null accuracy score: {0:0.4f}'.format(null_accuracy))
```

Null accuracy score: 0.7582

```
In [67]: from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion matrix

```
[[5999 1408]
 [ 465 1897]]
```

True Positives(TP) = 5999

True Negatives(TN) = 1897


False Positives(FP) = 1408

False Negatives(FN) = 465

```
In [68]: cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0']
                                index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[68]: <AxesSubplot:>



```
In [69]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.93	0.81	0.86	7407
>50K	0.57	0.80	0.67	2362
accuracy			0.81	9769
macro avg	0.75	0.81	0.77	9769

weighted avg	0.84	0.81	0.82	9769
--------------	------	------	------	------

```
In [70]: TP = cm[0,0]
         TN = cm[1,1]
         FP = cm[0,1]
         FN = cm[1,0]
```

```
In [71]: classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

         print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))

Classification accuracy : 0.8083
```

```
In [72]: classification_error = (FP + FN) / float(TP + TN + FP + FN)

         print('Classification error : {0:0.4f}'.format(classification_error))

Classification error : 0.1917
```

```
In [73]: precision = TP / float(TP + FP)

         print('Precision : {0:0.4f}'.format(precision))

Precision : 0.8099
```

```
In [74]: recall = TP / float(TP + FN)

         print('Recall or Sensitivity : {0:0.4f}'.format(recall))

Recall or Sensitivity : 0.9281
```

```
In [75]: true_positive_rate = TP / float(TP + FN)

         print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))

True Positive Rate : 0.9281
```

```
In [76]: false_positive_rate = FP / float(FP + TN)

         print('False Positive Rate : {0:0.4f}'.format(false_positive_rate))

False Positive Rate : 0.4260
```

```
In [77]: specificity = TN / (TN + FP)

         print('Specificity : {0:0.4f}'.format(specificity))

Specificity : 0.5740
```

```
In [78]: y_pred_prob = gnb.predict_proba(X_test)[0:10]

         y_pred_prob
```

```
Out[78]: array([[9.99999426e-01, 5.74152436e-07],
                [9.99687907e-01, 3.12093456e-04],
                [1.54405602e-01, 8.45594398e-01],
                [1.73624321e-04, 9.99826376e-01],
                [8.20121011e-09, 9.99999992e-01],
```

```
[8.76844580e-01, 1.23155420e-01],
[9.99999927e-01, 7.32876705e-08],
[9.99993460e-01, 6.53998797e-06],
[9.87738143e-01, 1.22618575e-02],
[9.99999996e-01, 4.01886317e-09]])
```

```
In [79]: y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Prob of - <=50K', 'Prob of
y_pred_prob_df
```

```
Out[79]:
```

	Prob of - <=50K	Prob of - >50K
0	9.999994e-01	5.741524e-07
1	9.996879e-01	3.120935e-04
2	1.544056e-01	8.455944e-01
3	1.736243e-04	9.998264e-01
4	8.201210e-09	1.000000e+00
5	8.768446e-01	1.231554e-01
6	9.999999e-01	7.328767e-08
7	9.999935e-01	6.539988e-06
8	9.877381e-01	1.226186e-02
9	1.000000e+00	4.018863e-09

```
In [80]: gnb.predict_proba(X_test)[0:10, 1]
```

```
Out[80]: array([5.74152436e-07, 3.12093456e-04, 8.45594398e-01, 9.99826376e-01,
9.99999992e-01, 1.23155420e-01, 7.32876705e-08, 6.53998797e-06,
1.22618575e-02, 4.01886317e-09])
```

```
In [81]: y_pred1 = gnb.predict_proba(X_test)[: , 1]
```

```
In [82]: # adjust the font size
plt.rcParams['font.size'] = 12


# plot histogram with 10 bins
plt.hist(y_pred1, bins = 10)

# set the title of predicted probabilities
plt.title('Histogram of predicted probabilities of salaries >50K')

# set the x-axis limit
plt.xlim(0,1)

# set the title
plt.xlabel('Predicted probabilities of salaries >50K')
plt.ylabel('Frequency')
```

```
Out[82]: Text(0, 0.5, 'Frequency')
```



In [83]:

```
# plot ROC Curve

from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred1, pos_label = '>50K')

plt.figure(figsize=(6,4))

plt.plot(fpr, tpr, linewidth=2)

plt.plot([0,1], [0,1], 'k--' )

plt.rcParams['font.size'] = 12


plt.title('ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries')

plt.xlabel('False Positive Rate (1 - Specificity)')

plt.ylabel('True Positive Rate (Sensitivity)')

plt.show()
```

ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries



In [84]:

```
# compute ROC AUC

from sklearn.metrics import roc_auc_score
```

```
ROC_AUC = roc_auc_score(y_test, y_pred1)

print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC AUC : 0.8941

```
In [85]: # calculate cross-validated ROC AUC

from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(gnb, X_train, y_train, cv=5, scoring='roc_
print('Cross validated ROC AUC : {:.4f}'.format(Cross_validated_ROC_AUC))
```

Cross validated ROC AUC : 0.8938

```
In [86]: # Applying 10-Fold Cross Validation

from sklearn.model_selection import cross_val_score

scores = cross_val_score(gnb, X_train, y_train, cv = 10, scoring='accuracy')

print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.81359649 0.80438596 0.81175954 0.8056165 0.79596314 0.79684072 0.81044318 0.81175954 0.80210619 0.81044318]

```
In [87]: # compute Average cross-validation score

print('Average cross-validation score: {:.4f}'.format(scores.mean()))
```

Average cross-validation score: 0.8063

In []: